
DISTRIBUTED CONFIGURATION WITH CLOJURE

AGENDA

- ▶ Definition
- ▶ Key Aspects
- ▶ Available Solutions
- ▶ Can we do better?
- ▶ Demo



WHO?

- ▶ @reborg
- ▶ <https://github.com/reborg>
- ▶ Hacking @ Mailonline - London
- ▶ "Clojure Weekly" @ <http://reborg.net>



DEFINITION

Configuration is any “external handle” that greatly affects the application behaviour without the need for re-writing it.

EXAMPLES

- ▶ environments (int, prod, test etc.)
- ▶ user based (e.g. account settings)
- ▶ multi-tenant applications (blurring into proper feature)
- ▶ feature toggles (per-request on/off checks)
- ▶ A/B tests (toggling between two features)

VALUE LIFETIME

- ▶ Bootstrap (no need for "if" statements)
- ▶ Scheduled time (e.g. daily)
- ▶ Real-time (requires condition checks)

STORAGE

- ▶ Sources (same as hosting language)
- ▶ External files (potentially dedicated format)
- ▶ Relational DB (easily accessible, real-time)
- ▶ Key-Value stores (separation from other application data)
- ▶ Dedicated service (just for config)

ORGANISATION

- ▶ Single file (or unit)
- ▶ Multiple files/units (might need merge for reading)
- ▶ Relational (tables)
- ▶ Hierarchical

CACHING

- ▶ Boost performance for reading configuration
- ▶ Needs to be invalidated on config changes
- ▶ Active check on timestamps
- ▶ TTL entries in the cache
- ▶ Pull: on-demand read, scheduled read
- ▶ Push: nohup, queues, custom sockets

DISTRIBUTION

- ▶ For clustered applications
- ▶ Changes need to be distributed across the cluster
- ▶ Changes need to happen roughly at the same time
- ▶ Local caches might need explicit invalidation

VERSIONING

- ▶ A must have for real-time configurations
- ▶ Need for tracking/auditing changes
- ▶ Revert when necessary
- ▶ “Diffing”: what is about to change, what was changed

KEY ASPECTS

COMPARING OPTIONS


	Plain text	MySQL	Redis	Etcd
Updating	Custom	Triggers	Keyspace notifications	Watchers
Versioning	Any VCS	Custom schema	custom time-series	custom time-series
Diffing	Easy	Custom	Custom	Custom
Organisation	Basic	Tabular	Hierarchical/Tabular	Hierarchical/Tabular
Distribution	Custom	Built-in	Built-in	Built-in
Format	Any	Binary/Proprietary	Binary/Proprietary	Binary/Proprietary
Dedicated?	Yes	No	No	No

CAN WE DO BETTER?

- ▶ Plain Text (versioning, diffs, any format, dedicated)
- ▶ Hierarchical model (flexible configuration)
- ▶ Push model (immediate reaction to changes)
- ▶ Dedicated to configuration only
- ▶ Distributed, Resilient, Scalable

SOLUTIONS

PROPOSED MODEL



WHY FUNCTIONAL?

- ▶ Immutable data structures perfect fit for versioning
- ▶ Built-in CAS semantic for local cache changes
- ▶ “Liveness” of the system allows instant re-evaluation
- ▶ Easy concurrency for distributed environments

CLOJURE

- ▶ Aleph/Manifold to deal with WebSockets (TCP or UDP)
- ▶ Java.NIO watch service for file changes notification
- ▶ Atoms for local caches and concurrent updates
- ▶ Metadata (stores information on raw-sources)
- ▶ Runs on any Java install

WHAT'S IN THE DEMO?

- ▶ 4 servers, 2-8 clients
- ▶ configuration folder: `fluorinedemo/apps`
- ▶ clients subscribe to `apps/test-json` subfolder
- ▶ configuration sent on client startup
- ▶ file-system as a tree, merged into hash-map
- ▶ will change one configuration parameter in one file

CURRENT SOLUTION TECHNICALITIES

- ▶ keep track of reconnecting clients
- ▶ keep-alive pings
- ▶ client local caches
- ▶ supporting edn and json
- ▶ JavaScript (and others) Client: WIP
- ▶ code and demo: <https://github.com/reborg/fluorine>

~ FIN ~

PDF SLIDES AVAILABLE: [HTTP://TINYURL.COM/ZYAE7U3](http://tinyurl.com/zyae7u3)